Shangtong Zhang

Department of Computer Science
School of Engineering and Applied Science
University of Virginia
Charlottesville, VA, 22903

Rice Hall 422
85 Engineer's Way
Charlottesville, VA, 22903
shangtong@virginia.edu

RESEARCH INTEREST

The goal of my research is to solve sequential decision making problems in a scalable and reliable way. Currently, I focus on Reinforcement Learning (RL) as a solution method. In particular, I work on stochastic approximations for RL, theories and algorithms of RL, and applications by RL.

ACADEMIC EMPLOYMENTS

Assistant Professor   Aug 2022 - Present
Department of Computer Science
University of Virginia, VA, United States

Research Scientist Interns
Microsoft Research Montreal   Jun 2021 - Sep 2021
DeepMind London   Feb 2021 - Jun 2021
Microsoft Research Montreal   Jun 2020 - Aug 2020

EDUCATION

Doctor of Philosophy, Computer Science   Oct 2018 - Jul 2022
University of Oxford, Oxford, United Kingdom
Advisor: Prof. Shimon Whiteson
Thesis: Breaking the Deadly Triad in Reinforcement Learning

Master of Science, Computer Science   Sep 2016 - Jul 2018
University of Alberta, Edmonton, Canada
Advisor: Prof. Richard S. Sutton
Thesis: Learning with Artificial Neural Networks

Bachelor of Science, Computing Science   Sep 2012 - Jul 2016
Fudan University, Shanghai, China

PUBLICATIONS

Advisees of SZ are underlined; * indicates equal contribution; † indicates equal supervision.

Preprints

(P1) *Experience Replay Addresses Loss of Plasticity in Continual Learning.*
Jiuqi Wang, Rohan Chandra, **Shangtong Zhang**.

ArXiv Preprint, 2025.

(P2) *GameChat: Multi-LLM Dialogue for Safe, Agile, and Socially Optimal Multi-Agent Navigation in Constrained Environments.*
Vagul Mahadevan, **Shangtong Zhang**, Rohan Chandra.
ArXiv Preprint, 2025.

(P3) *Towards Large Language Models that Benefit for All: Benchmarking Group Fairness in Reward Models.*
Kefan Song, Jin Yao, Runan Jiang, Rohan Chandra, **Shangtong Zhang**.
ArXiv Preprint, 2025.

(P4) *Group Fairness in Multi-Task Reinforcement Learning.*
Kefan Song, Runan Jiang, Rohan Chandra, **Shangtong Zhang**.
ArXiv Preprint, 2025.

(P5) *A Survey of In-Context Reinforcement Learning.*
Amir Moeini, Jiuqi Wang, Jakob Beck, Ethan Blaser, Shimon Whiteson, Rohan Chandra, **Shangtong Zhang**.
ArXiv Preprint, 2025.

(P6) *CRASH: Challenging Reinforcement-Learning Based Adversarial Scenarios For Safety Hardening.*
Amar Kulkarni, **Shangtong Zhang**, Madhur Behl.
ArXiv Preprint, 2024.

(P7) *Almost Sure Convergence Rates and Concentration of Stochastic Approximation and Reinforcement Learning with Markovian Noise.*
Xiaochi Qian*, Xinyu Liu*, Zixuan Xie*, **Shangtong Zhang**.
ArXiv Preprint, 2024.

(P8) *Asymptotic and Finite Sample Analysis of Nonexpansive Stochastic Approximations with Markovian Noise.*
Ethan Blaser, **Shangtong Zhang**.
ArXiv Preprint, 2024.

(P9) *Almost Sure Convergence of Linear Temporal Difference Learning with Arbitrary Features.*
Jiuqi Wang, **Shangtong Zhang**.
ArXiv Preprint, 2024.

(P10) *AlphaStar Unplugged: Large Scale Offline Reinforcement Learning.*
Michael Mathieu*, Sherjil Ozair*, Srivatsan Srinivasan*, Caglar Gulcehre*, **Shangtong Zhang***, Ray Jiang*, Tom Le Paine*, Richard Powell, Konrad Zolna, Julian Schrittwieser, David Choi, Petko Georgiev, Daniel Kenji Toyama, Aja Huang, Roman Ring, Igor Babuschkin, Timo Ewalds, Mahyar Bordbar, Sarah Henderson, Sergio Gomez Colmenarejo, Aaron van den Oord, Wojciech M. Czarnecki, Nando de Freitas, Oriol Vinyals.
ArXiv Preprint, 2023

Invited Articles

(I1) *A New Challenge in Policy Evaluation.*
Shangtong Zhang.
AAAI Conference on Artificial Intelligence (**AAAI**), 2023
New Faculty Highlights Program.

Refereed Journals

(J1) *The ODE Method for Stochastic Approximation and Reinforcement Learning with Markovian Noise.*
Shuze Liu, Shuhang Chen, **Shangtong Zhang**.
Journal of Machine Learning Research (**JMLR**), 2025.

(J2) *Global Optimality and Finite Sample Analysis of Softmax Off-Policy Actor Critic under State Distribution Mismatch.*
Shangtong Zhang, Remi Tachet des Combes[‡], Romain Laroche[‡].
Journal of Machine Learning Research (**JMLR**), 2022.

(J3) *Truncated Emphatic Temporal Difference Methods for Prediction and Control.*
Shangtong Zhang, Shimon Whiteson.
Journal of Machine Learning Research (**JMLR**), 2022.

(J4) *MLPack 3: A Fast, Flexible Machine Learning Library.*
Ryan Curtin, Marcus Edel, Mikhail Lozhnikov, Yannis Mentekidis, Sumedh Ghaisas, **Shangtong Zhang**
Journal of Open Source Software (**JOSS**), 2018.

Refereed Conference Papers

(C1) *Linear Q-Learning Does Not Diverge in L^2: Convergence Rates to a Bounded Set.*
Xinyu Liu[*], Zixuan Xie[*], **Shangtong Zhang**.
International Conference on Machine Learning (**ICML**), 2025.
Acceptance rate: 26.9%

(C2) *Transformers Can Learn Temporal Difference Methods for In-Context Reinforcement Learning.*
Jiuqi Wang[*], Ethan Blaser[*], Hadi Daneshmand, **Shangtong Zhang**.
International Conference on Learning Representations (**ICLR**), 2025.
Acceptance rate: 32.08%.
QuantCo Spotlight Award at the ICML Workshop on In-Context Learning, 2024.

(C3) *Revisiting a Design Choice in Gradient Temporal Difference Learning.*
Xiaochi Qian, **Shangtong Zhang**.
International Conference on Learning Representations (**ICLR**), 2025.
Acceptance rate: 32.08%.

(C4) *Efficient Policy Evaluation with Safety Constraint for Reinforcement Learning.*
Claire Chen[*], Shuze Liu[*], **Shangtong Zhang**.
International Conference on Learning Representations (**ICLR**), 2025.
Acceptance rate: 32.08%.

(C5) *Doubly Optimal Policy Evaluation for Reinforcement Learning.*
Shuze Liu, Claire Chen, **Shangtong Zhang**.
International Conference on Learning Representations (**ICLR**), 2025.
Acceptance rate: 32.08%.

(C6) *Efficient Multi-Policy Evaluation for Reinforcement Learning.*
Shuze Liu, Claire Chen, **Shangtong Zhang**.
AAAI Conference on Artificial Intelligence (**AAAI**), 2025.
Acceptance rate: 23.4%. **Oral presentation (4.6%)**.

(C7) *Efficient Policy Evaluation with Offline Data Informed Behavior Policy Design.*
Shuze Liu, **Shangtong Zhang**.
International Conference on Machine Learning (**ICML**), 2024.
Acceptance rate: 27.5%

(C8) *On the Convergence of SARSA with Linear Function Approximation.*
Shangtong Zhang, Remi Tachet des Combes, Romain Laroche.
International Conference on Machine Learning (**ICML**), 2023.
Acceptance rate: 28%

(C9) *A Deeper Look at Discounting Mismatch in Actor-Critic Algorithms.*
Shangtong Zhang, Romain Laroche, Harm van Seijen, Shimon Whiteson, Remi Tachet des Combes.
International Conference on Autonomous Agents and Multiagent Systems (**AAMAS**), 2022.
Acceptance rate: 26%. **Oral presentation.**

(C10) *Learning Expected Emphatic Traces for Deep RL.*
Ray Jiang, **Shangtong Zhang**, Veronica Chelu, Adam White, Hado van Hasselt.
AAAI Conference on Artificial Intelligence (**AAAI**), 2022.
Acceptance rate: 15%.

(C11) *Breaking the Deadly Triad with a Target Network.*
Shangtong Zhang, Hengshuai Yao, Shimon Whiteson.
International Conference on Machine Learning (**ICML**), 2021.
Acceptance rate: 21.5%.

(C12) *Average-Reward Off-Policy Policy Evaluation with Function Approximation.*
Shangtong Zhang*, Yi Wan*, Richard S. Sutton, Shimon Whiteson.
International Conference on Machine Learning (**ICML**), 2021.
Acceptance rate: 21.5%.

(C13) *Mean-Variance Policy Iteration for Risk-Averse Reinforcement Learning.*
Shangtong Zhang, Bo Liu, Shimon Whiteson.
AAAI Conference on Artificial Intelligence (**AAAI**), 2021.
Acceptance rate: 21.4%.

(C14) *Learning Retrospective Knowledge with Reverse Reinforcement Learning.*
Shangtong Zhang, Vivek Veeriah, Shimon Whiteson.
Conference on Neural Information Processing Systems (**NeurIPS**), 2020.
Acceptance rate: 20.1%.

(C15) *GradientDICE: Rethinking Generalized Offline Estimation of Stationary Values.*
Shangtong Zhang, Bo Liu, Shimon Whiteson.
International Conference on Machine Learning (**ICML**), 2020.
Acceptance rate: 21.8%.

(C16) *Provably Convergent Two-Timescale Off-Policy Actor-Critic with Function Approximation.*
Shangtong Zhang, Bo Liu, Hengshuai Yao, Shimon Whiteson.
International Conference on Machine Learning (**ICML**), 2020.
Acceptance rate: 21.8%.

(C17) *Deep Residual Reinforcement Learning.*
Shangtong Zhang, Wendelin Boehmer, Shimon Whiteson.
International Conference on Autonomous Agents and Multiagent Systems (**AAMAS**), 2020.
Acceptance rate: 23%. **Best Paper Award**.

(C18) *Mega-Reward: Achieving Human-Level Play without Extrinsic Rewards.*
Yuhang Song, Jianyi Wang, Thomas Lukasiewicz, Zhenghua Xu, **Shangtong Zhang**, Andrzej Wojcicki, Mai Xu.
AAAI Conference on Artificial Intelligence (**AAAI**), 2020.
Acceptance rate: 20.6%.

(C19) *DAC: The Double Actor-Critic Architecture for Learning Options.*
Shangtong Zhang, Shimon Whiteson.
Conference on Neural Information Processing Systems (**NeurIPS**), 2019.
Acceptance rate: 21.2%.

(C20) *Generalized Off-Policy Actor-Critic.*
Shangtong Zhang, Wendelin Boehmer, Shimon Whiteson.
Conference on Neural Information Processing Systems (**NeurIPS**), 2019.
Acceptance rate: 21.2%.

(C21) *Distributional Reinforcement Learning for Efficient Exploration.*
Borislav Mavrin, **Shangtong Zhang**, Hengshuai Yao, Linglong Kong, Kaiwen Wu, Yaoliang Yu
International Conference on Machine Learning (**ICML**), 2019.
Acceptance rate: 22.6%.

(C22) *ACE: An Actor Ensemble Algorithm for Continuous Control with Tree Search.*
Shangtong Zhang, Hao Chen, Hengshuai Yao.
AAAI Conference on Artificial Intelligence (**AAAI**), 2019.
Acceptance rate: 16.2%. **Spotlight presentation.**

(C23) *QUOTA: The Quantile Option Architecture for Reinforcement Learning.*
Shangtong Zhang, Borislav Mavrin, Linglong Kong, Bo Liu, Hengshuai Yao.
AAAI Conference on Artificial Intelligence (**AAAI**), 2019.
Acceptance rate: 16.2%. **Oral presentation.**

(C24) *Crossprop: Learning Representations by Stochastic Meta-Gradient Descent in Neural Networks.*
Vivek Veeriah*, **Shangtong Zhang***, Richard S. Sutton.
European Conference on Machine Learning and Principles and Practice of Knowledge Discovery in
Databases (**ECML-PKDD**), 2017.
Acceptance rate: 27.1%.

(C25) *A Deep Neural Network for Modeling Music.*
Pengjing Zhang, Xiaoqing Zheng, Wenqiang Zhang, Siyan Li, Sheng Qian,
Wenqi He, **Shangtong Zhang**, Ziyuan Wang
International Conference on Multimedia Retrieval (**ICMR**), 2015.
Acceptance rate: 31%.

Refereed Workshop Papers (Non-Archival)

(W1) *A Deeper Look at Experience Replay.*
Shangtong Zhang, Richard S. Sutton.
Deep RL Symposium at NIPS, 2017.

(W2) *Comparing Deep Reinforcement Learning and Evolutionary Methods in Continuous Control.*
Shangtong Zhang, Osmar R. Zaiane
Deep RL Symposium at NIPS, 2017

(W3) *A Demon Control Architecture with Off-Policy Learning and Flexible Behavior Policy.*
Shangtong Zhang, Richard S. Sutton.
Hierarchical RL Workshop at NIPS, 2017.

GRANTS

Enhancing the Security of Large Language Models Against Persuasion-Based Jailbreak Attacks
in Multi-Turn Dialogues
CCI Coastal Virginia Node, **Co-PI**, Total $60,000, My Share $20,000   2025 - 2026
RAMPART: Reinforcement Against Malicious Penetration by Adversaries in Realistic Topologies.
DARPA HR001123S0002, **Co-PI**, Total $2,150,000, My Share $77,000.   2023 - 2027
SLES: CRASH: Challenging Reinforcement-Learning Based Adversarial Scenarios for Safety Hardening.
NSF 2331904, **Co-PI**, Total $800,000, My Share $400,000.   2023 - 2026
III: Small: Moving Offline Learning to Rank Online, from Theory to Practice.
NSF 2128019, **PI**, Total $500,000, My Share $500,000.   2021 - 2025

HONORS

QuantCo Spotlight Award at the ICML In-Context Learning Workshop, 2024
Rising Stars in AI by KAUST AI Initiative, 2024
AAAI New Faculty Highlights, 2023

SERVICES

Organizers
CPS Rising Star Workshop 2024, Co-Chair

Panelists
National Science Foundation (NSF), 2024
Virginia's Commonwealth Cyber Initiative (CCI), 2024

Meta Meta Reviewer
RL Conference 2025 (Senior Area Chair)

Meta Reviewer
NeurIPS 2025 (Area Chair)
ICML 2025 (Area Chair)
ICLR 2024, 2025 (Area Chair)
L4DC 2025 (Program Committee)
AAMAS 2025 (Senior Program Committee)
RL Conference 2024 (Area Chair)
AISTATS 2024, 2025 (Area Chair)
ACML 2022, 2023, 2024 (Area Chair)

Reviewer
Transactions on Pattern Analysis and Machine Intelligence (1)
Transaction of Machine Learning Research (2)
Journal of Machine Learning Research (3)
Artificial Intelligence Journal (2)
Transactions on Intelligent Systems and Technology (2)
IJCAI 2023
AISTATS 2022
NeurIPS 2020, 2021, 2022, 2023
ICML 2020, 2021, 2022, 2023
AAAI 2020, 2021, 2022, 2023
ICLR 2021, 2022, 2023
SIGCOMM 2022
Offline Reinforcement Learning Workshop at NeurIPS 2020, 2021, 2022
Deep Reinforcement Learning Workshop at NeurIPS 2019, 2020, 2021, 2022
Adaptive and Learning Agents Workshop at AAMAS 2019, 2020
Optimization Foundations for Reinforcement Learning Workshop at NeurIPS 2019
Reinforcement Learning for Real Life Workshop at ICML 2019, 2021
Reinforcement Learning for Real Life Workshop at NeurIPS 2022

Conference Session Chair
AAAI 2023, "Reinforcement Learning Theory & Algorithms"

SUPERVISION

Doctral Students

Shuze Liu	2022 - Now
Ethan Blaser, NSF Graduate Research Fellowship	2023 - Now
Jiuqi Wang	2023 - Now
Xinyu Liu	2024 - Now
Zixuan Xie, UVA Provost's Fellowship, UVA Engineering Distinguished Fellowship	2024 - Now
Amir Moeini	2024 - Now

Alumni with Theses or Publications

Kefan Song, Master, UVA	2023 - 2024
Licheng Luo, Master, UVA	2023 - 2024
Steve Zhou, Undergraduate, UVA, DMP thesis	2023 - 2024
Xiaochi Qian, Undergraduate, University of Oxford, Research Intern	2022 - 2024

PhD Committees
Tao Jin (UVA, by Prof. Farzad Farnoud)
Fan Yao (UVA, by Prof. Hongning Wang and Prof. Haifeng Xu)
Ingy ElSayed-Aly (UVA, by Prof. Lu Feng)
Sudhir Shenoy (UVA, by Prof. Afsaneh Doryab)
Chuanhao Li (UVA, by Prof. Hongning Wang)

INVITED TALKS

Understanding the Training and Inference of Reinforcement Learning

Tsinghua University, hosted by Prof. Hongning Wang	June 2024

On the Cheating of Offline Reinforcement Learning

KAUST Rising Stars in AI Symposium	Feb 2024

Offline Reinforcement Learning: Current and Future

AAAI New Faculty Highlight Program	Feb 2023

Breaking the Deadly Triad in Off-Policy Reinforcement Learning

Department of Computer Science, University of Virginia	Mar 2022
School of Computing Science, Simon Fraser University	Feb 2022
Department of Electrical & Computer Engineering, University of Waterloo	Feb 2022
School of Informatics, University of Edinburgh	Oct 2021

Breaking the Deadly Triad with a Target Network

Microsoft Research Summit	Oct 2021

Breaking the Deadly Triad in Reinforcement Learning

RL team, DeepMind, hosted by Dr. Hado van Hasselt	Sep 2021

Off-Policy Evaluation

Data Fest 2020, Open Data Science	Oct 2020

Off-Policy Evaluation and Control

ByteDance AI Lab, Shanghai	Oct 2020

Coding Deep RL Papers

NIPS MLTrain Workshop, Long Beach	Dec 2019

Off-Policy Actor-Critic Algorithms

Latent Logic LTD, Oxford	Apr 2019

<div align="center">**TEACHING**</div>

University of Virginia

CS 4501: Reinforcement Learning	Fall 2024
CS 6316: Machine Learning	Spring 2024
CS 4501: Optimization	Fall 2023
CS 6501: Topics in Reinforcement Learning	Fall 2022

<div align="center">**OPEN SOURCE CODE**</div>

GitHub Repo: PyTorch Deep RL
A zoo of popular deep RL algorithms in PyTorch with **3k stars**.

GitHub Repo: Reinforcement Learning: An Introduction
Python implementation of the book *Reinforcement Learning: An Introduction* with **13.8k stars**.

Google Summer of Code (GSoC)

MLPack	2017
The Xapian Project	2014