

Shangtong Zhang

422 Rice Hall, 85 Engineer's Way, Charlottesville, VA, 22903, United States
shangtong@virginia.edu <https://shangtongzhang.github.io>

Research Interest	The goal of my research is to solve sequential decision making problems in a scalable and reliable way. Currently, I focus on reinforcement learning as a solution method.	
Academic Employments	Assistant Professor Department of Computer Science University of Virginia, United States	Aug. 2022 - Present
Education	University of Oxford , United Kingdom Doctor of Philosophy in Computer Science Advisor: Prof. Shimon Whiteson	Oct. 2018 - July. 2022
	University of Alberta , Canada Master of Science in Computer Science, Advisor: Prof. Richard S. Sutton	Sept. 2016 - Aug. 2018
	Fudan University , China Bachelor of Science in Computing Science	Sept. 2012 - Jun. 2016
Publications	<ol style="list-style-type: none"><u>Direct Gradient Temporal Difference Learning</u> Xiaochi Qian^[advisee], Shangtong Zhang. arXiv:2308.01170, 2023.<u>Improving Monte Carlo Evaluation with Offline Data</u> Shuze Liu^[advisee], Shangtong Zhang. arXiv:2301.13734, 2023.<u>On the Convergence of SARSA with Linear Function Approximation</u> Shangtong Zhang, Remi Tachet des Combes, Romain Laroche. International Conference on Machine Learning (ICML), 2023. Acceptance rate: 28%<u>Global Optimality and Finite Sample Analysis of Softmax Off-Policy Actor Critic under State Distribution Mismatch</u> Shangtong Zhang, Remi Tachet des Combes[†], Romain Laroche[‡]. Journal of Machine Learning Research (JMLR), 2022.<u>Truncated Emphatic Temporal Difference Methods for Prediction and Control</u> Shangtong Zhang, Shimon Whiteson. Journal of Machine Learning Research (JMLR), 2022.<u>A Deeper Look at Discounting Mismatch in Actor-Critic Algorithms</u> Shangtong Zhang, Romain Laroche, Harm van Seijen, Shimon Whiteson, Remi Tachet des Combes. International Conference on Autonomous Agents and Multiagent Systems (AAMAS), 2022. Acceptance rate: 26%<u>Learning Expected Emphatic Traces for Deep RL</u> Ray Jiang, Shangtong Zhang, Veronica Chelu, Adam White, Hado van Hasselt.	

AAAI Conference on Artificial Intelligence (**AAAI**), 2022.
Acceptance rate: 15%.

8. StarCraft II Unplugged: Large Scale Offline Reinforcement Learning
Michael Mathieu*, Sherjil Ozair*, Srivatsan Srinivasan*, Caglar Gulcehre*,
Shangdong Zhang*, Ray Jiang*, Tom Le Paine*, Richard Powell, Konrad Zolna,
Julian Schrittwieser, David Choi, Petko Georgiev, Daniel Kenji Toyama,
Aja Huang, Roman Ring, Igor Babuschkin, Timo Ewalds, Mahyar Bordbar,
Sarah Henderson, Sergio Gomez Colmenarejo, Aaron van den Oord,
Wojciech M. Czarnecki, Nando de Freitas, Oriol Vinyals.
arXiv:2308.03526, 2023
Deep RL Workshop at NeurIPS, 2021
9. Breaking the Deadly Triad with a Target Network
Shangdong Zhang, Hengshuai Yao, Shimon Whiteson.
International Conference on Machine Learning (**ICML**), 2021.
Acceptance rate: 21.5%.
10. Average-Reward Off-Policy Policy Evaluation with Function Approximation
Shangdong Zhang*, Yi Wan*, Richard S. Sutton, Shimon Whiteson.
International Conference on Machine Learning (**ICML**), 2021.
Acceptance rate: 21.5%.
11. Mean-Variance Policy Iteration for Risk-Averse Reinforcement Learning
Shangdong Zhang, Bo Liu, Shimon Whiteson.
AAAI Conference on Artificial Intelligence (**AAAI**), 2021.
Acceptance rate: 21.4%.
12. Learning Retrospective Knowledge with Reverse Reinforcement Learning
Shangdong Zhang, Vivek Veeriah, Shimon Whiteson.
Conference on Neural Information Processing Systems (**NeurIPS**), 2020.
Acceptance rate: 20.1%.
13. GradientDICE: Rethinking Generalized Offline Estimation of Stationary Values
Shangdong Zhang, Bo Liu, Shimon Whiteson.
International Conference on Machine Learning (**ICML**), 2020.
Acceptance rate: 21.8%.
14. Provably Convergent Two-Timescale Off-Policy Actor-Critic with Function Approximation
Shangdong Zhang, Bo Liu, Hengshuai Yao, Shimon Whiteson.
International Conference on Machine Learning (**ICML**), 2020.
Acceptance rate: 21.8%.
15. Deep Residual Reinforcement Learning
Shangdong Zhang, Wendelin Boehmer, Shimon Whiteson.
International Conference on Autonomous Agents and Multiagent Systems (**AAMAS**), 2020.
Acceptance rate: 23%.
Best Paper Award.
16. Mega-Reward: Achieving Human-Level Play without Extrinsic Rewards
Yuhang Song, Jianyi Wang, Thomas Lukasiewicz, Zhenghua Xu,
Shangdong Zhang, Andrzej Wojcicki, Mai Xu
AAAI Conference on Artificial Intelligence (**AAAI**), 2020.
Acceptance rate: 20.6%.
17. DAC: The Double Actor-Critic Architecture for Learning Options
Shangdong Zhang, Shimon Whiteson.
Conference on Neural Information Processing Systems (**NeurIPS**), 2019.
Acceptance rate: 21.2%.

18. Generalized Off-Policy Actor-Critic
Shangdong Zhang, Wendelin Boehmer, Shimon Whiteson.
 Conference on Neural Information Processing Systems (**NeurIPS**), 2019.
 Acceptance rate: 21.2%.
19. Distributional Reinforcement Learning for Efficient Exploration
 Borislav Mavrin, **Shangdong Zhang**[†], Hengshuai Yao, Linglong Kong,
 Kaiwen Wu, Yaoliang Yu
 International Conference on Machine Learning (**ICML**), 2019.
 Acceptance rate: 22.6%.
 A short version is accepted as an extended abstract at AAMAS 2019.
20. ACE: An Actor Ensemble Algorithm for Continuous Control with Tree Search
Shangdong Zhang, Hao Chen, Hengshuai Yao.
 AAAI Conference on Artificial Intelligence (**AAAI**), 2019.
 Acceptance rate: 16.2%.
21. QUOTA: The Quantile Option Architecture for Reinforcement Learning
Shangdong Zhang, Borislav Mavrin, Linglong Kong, Bo Liu, Hengshuai Yao.
 AAAI Conference on Artificial Intelligence (**AAAI**), 2019.
 Acceptance rate: 16.2%.
22. MLPack 3: A Fast, Flexible Machine Learning Library
 Ryan Curtin, Marcus Edel, Mikhail Lozhnikov, Yannis Mentekidis, Sumedh Ghaisas,
Shangdong Zhang
 Journal of Open Source Software (**JOSS**), 2018.
23. Crossprop: Learning Representations by Stochastic Meta-Gradient Descent
in Neural Networks
 Vivek Veeriah*, **Shangdong Zhang***, Richard S. Sutton.
 European Conference on Machine Learning and Principles and Practice of Knowledge
 Discovery in Databases (**ECML-PKDD**), 2017.
 Acceptance rate: 27.1%.
24. A Deeper Look at Experience Replay
Shangdong Zhang, Richard S. Sutton.
Deep RL Symposium at NIPS, 2017.
25. Comparing Deep Reinforcement Learning and Evolutionary Methods
in Continuous Control
Shangdong Zhang, Osmar R. Zaiane
Deep RL Symposium at NIPS, 2017.
26. A Demon Control Architecture with Off-Policy Learning and Flexible Behavior
Policy
Shangdong Zhang, Richard S. Sutton.
Hierarchical RL Workshop at NIPS, 2017.
27. A Deep Neural Network for Modeling Music
 Pengjing Zhang, Xiaoqing Zheng, Wenqiang Zhang, Siyan Li, Sheng Qian,
 Wenqi He, **Shangdong Zhang**, Ziyuan Wang
 International Conference on Multimedia Retrieval (**ICMR**), 2015.
 Acceptance rate: 31%.

*: Equal contribution

[‡]: Equal advising

[†] : My name does not appear in the ICML proceedings due to a mistake in submission.
 See Acknowledgments, arXiv, or AAMAS proceedings for clarification.

Services	Meta Reviewer & Area Chair ACML 2022, 2023	
	Expert Reviewer ICML 2021	
	Reviewer & Program Committee Transactions on Pattern Analysis and Machine Intelligence (1) Transaction of Machine Learning Research (2) Journal of Machine Learning Research (3) Artificial Intelligence Journal (2) Transactions on Intelligent Systems and Technology (2) IJCAI 2023 AISTATS 2022 NeurIPS 2020, 2021, 2022, 2023 ICML 2020, 2022, 2023 AAAI 2020, 2021, 2022, 2023 ICLR 2021, 2022, 2023 SIGCOMM 2022 Offline Reinforcement Learning Workshop at NeurIPS 2020, 2021, 2022 Deep Reinforcement Learning Workshop at NeurIPS 2019, 2020, 2021, 2022 Adaptive and Learning Agents Workshop at AAMAS 2019, 2020 Optimization Foundations for Reinforcement Learning Workshop at NeurIPS 2019 Reinforcement Learning for Real Life Workshop at ICML 2019, 2021 Reinforcement Learning for Real Life Workshop at NeurIPS 2022	
	Session Chairs AAAI 2023, “Reinforcement Learning Theory & Algorithms”	
	Departmental Services Graduate Admission Committee at CS of UVA, 2022 - 2023 AY	
Honours	<i>Runner-Up for the IFAAMAS Victor Lesser Dissertation Award</i>	2022
	<i>Alf Weaver Junior Faculty Fellowship, University of Virginia</i>	2022 - 2027
	<i>EPSRC Studentship, University of Oxford</i>	2018 - 2022
	<i>AAMAS Student Scholarship</i>	2022
	<i>ICLR Outstanding Reviewer</i>	2021
	<i>NeurIPS Reviewer Award</i>	2020
	<i>ICML Reviewer Award</i>	2020
	<i>Light Senior Scholarship, St Catherine’s College, University of Oxford</i>	2020
	<i>AAMAS Travel Award</i>	2020
	<i>AAMAS Best Paper Award</i>	2020
	<i>NeurIPS Optimization Foundations for RL Workshop Travel Award</i>	2019
	<i>NeurIPS Travel Award</i>	2019
	<i>AAAI Travel Award</i>	2019
	<i>NIPS Hierarchical RL Workshop Travel Award</i>	2017
	<i>EMC Scholarship, Fudan University</i>	2014
Supervision	PhD Students Shuze Liu (2022 -) Jiuqi Wang (2023 -)	

Undergraduates

Ja-Zhua Cheng (2022 -)

Research Assistants

Xiaochi Qian (2022 -)

Alumni

Zhengkun Xiao (Master, 2022 - 2023, now PhD student at University of Florida)

Jiuqi Wang (RA, 2022 - 2023, now PhD student at University of Virginia)

PhD Committees

Ingy ElSayed-Aly (by Prof. Lu Feng at UVA)

Sudhir Shenoy (by Prof. Afsaneh Doryab at UVA)

Chuanhao Li (by Prof. Hongning Wang at UVA)

Qualification:

Zeyu Mu (by Prof. Brian Park at UVA)

Dane Williamson (by Prof. Yangfeng Ji at UVA)

Matthew Landers (by Prof. Afsaneh Doryab at UVA)

Ethan Harrison Blaser (by Prof. Hongning Wang at UVA)

Invited Talks	<i>Offline Reinforcement Learning: Current and Future</i>	2023
	AAAI New Faculty Highlight Program	
	<i>Breaking the Deadly Triad in Off-Policy Reinforcement Learning</i>	
	School of Computing Science, Simon Fraser University	2022
	Department of Electrical & Computer Engineering, University of Waterloo	2022
	Department of Computer Science, University of Virginia	2022
	School of Informatics, University of Edinburgh	2021
	<i>Breaking the Deadly Triad in Reinforcement Learning</i>	2021
	RL team, DeepMind, hosted by Hado van Hasselt	
	<i>Breaking the Deadly Triad with a Target Network</i>	2021
	Microsoft Research Summit	
	<i>Off-Policy Evaluation</i>	2020
	Data Fest 2020, Open Data Science	
	<i>Off-Policy Evaluation and Control</i>	2020
	ByteDance AI Lab, Shanghai	
	<i>Off-Policy Actor-Critic Algorithms</i>	2019
	Latent Logic LTD, Oxford	
	<i>Generalized Off-Policy Actor-Critic</i>	2019
	Noah's Ark Lab, Huawei, Edmonton	
	<i>Exploration with Quantile Options</i>	2018
	Huawei RL Workshop, Edmonton	
	<i>Coding Deep RL Papers</i>	2017
	NIPS MLTrain Workshop, Long Beach	

Teaching	University of Virginia, Instructor CS6501: Topics in Reinforcement Learning	Fall 2022
	University of Oxford, Teaching Assistant AIMS CDT Lectures	Michaelmas 2019
	University of Alberta, Teaching Assistant CMPUT 229 Computer Organization and Architecture	Fall 2016
Research Internships	Microsoft Research Montreal, Canada Collaboration: Remi Tachet des Combes, Romain Laroche, and Harm van Seijen	Jun. 2021 - Sept. 2021
	DeepMind London, United Kingdom Collaboration: AlphaStar team (Michael Mathieu, Oriol Vinyals, etc) Collaboration: Adam White and Hado van Hasselt	Feb. 2021 - Jun. 2021
	DeepDrive, Edmonton, Canada Collaboration: Hengshuai Yao	Sept. 2020 - Dec. 2020
	Microsoft Research Montreal, Canada Collaboration: Remi Tachet des Combes, Romain Laroche, and Harm van Seijen	Jun. 2020 - Aug. 2020
	Noah's Ark Lab, Huawei, Edmonton, Canada Collaboration: Hengshuai Yao	May. 2018 - Aug. 2018
Code	PyTorch Deep RL A zoo of popular deep RL algorithms in PyTorch with 2.8k stars in Github.	
	Reinforcement Learning: An Introduction Python implementation of the book <i>Reinforcement Learning: An Introduction</i> with 11.7k stars in Github.	
	Google Summer of Code (GSoC) 2017 Contributed to MLPack by implementing a deep RL framework.	
	Google Summer of Code (GSoC) 2014 Contributed to Xapian by optimizing the post list and the position list.	